

A1 – CUDA Introduction

Bidyadhar Mohanty(16CO212)
Soham Patil(16CO249)

Q.2. 1.How many floating operations are being performed in the matrix addition kernel?

Number of floating operations performed are same as the size of the matrices i.e.
 $\text{row_size} * \text{column_size}$.

2. How many global memory reads are being performed by your kernel?

For every thread, two global reads are performed, one for `host_mat1` and the other for `host_mat2`.

3. How many global memory writes are being performed by your kernel?

Only one global write takes place per thread to `host_sum`.